

COL729 Assignment1 Report

Manish Yadav
2015CS10460

1 LLVM IR bitcode

The -O0 doesn't consist of any optimization and hence has a corresponding instruction for each instruction in source code. Loop is implemented using different modules each having a branch instruction at the end. For every variable (including argc and argv), space is allocated on the stack first using alloca instruction, the value is stored in the stack using store and when we need the variable, it is loaded in a temporary variable %n using load (n is an integer).

In -O2, the compiler can detect empty loops and completely omit them in IR. The variables are used directly instead of first storing them in stack (e.g. `- icmp sgt i32 %argc, 1` uses argc directly). Some instructions are also changed. For instance, `argc >= 2` uses a `sgt` instruction instead of `sge` which is used in -O0. In cases where tail call optimizations can be done, it is usually done. It can also detect loops like `sumn()` which sums the first n integers and use the closed formula to compute the sum instead of iterating in the loop. phi instruction is used in many places.

1.1 emptyloop

1.1.1 -O0

It creates space for variables on stack and stores the values in it initially using alloca and store instructions.

```
define i32 @emptyloop(i32 %argc, i8** %argv) #0 {  
    %1 = alloca i32, align 4  
    %2 = alloca i8**, align 8  
    %i = alloca i64, align 8  
    %numiter = alloca i64, align 8  
    store i32 %argc, i32* %1, align 4  
    store i8** %argv, i8*** %2, align 8
```

```

store i64 2147483646, i64* %numiter, align 8
%3 = load i32* %1, align 4
%4 = icmp sge i32 %3, 2
br i1 %4, label %5, label %11

```

Temporary variables %1 and %2 are created to hold the values for argc and argv(pointer) in stack. %3 is created to hold the temporary value which is loaded from stack(it too contains argc). %4 = icmp sge i32 %3, 2 is the corresponding instruction for argc >= 2 used in if statement. %numiter stores 2147483646 which is the value of INT_MAX - 1 indicating that it directly computes the value without using a separate sub instruction for it. Depending on the result, it branches off to labels %5 or %11.

```

; <label>:5                                     ; preds = %0
%6 = load i8*** %2, align 8
%7 = getelementptr inbounds i8** %6, i64 1
%8 = load i8** %7, align 8
%9 = call i32 @atoi(i8* %8) #2
%10 = sext i32 %9 to i64
store i64 %10, i64* %numiter, align 8
br label %11

```

The above code shows bitcode for label %5. It executes the statement `numiter = atoi(argv[1])`. First the vales of argv (i.e. the address of the first element) is loaded in %6. `getelemenetptr` instruction is used to find the *address* of the 1st element of the array which is stored in %7. %8 then stores the value of argv[1] using that address. The `atoi()` function is called on argv[1] and the result is stored in %9. `sext (sign extends)` typecasts the result from i32 to i64 which is finally stored in %numiter. Then it branches off to label %11 which executes the assignment statement of for loop.

```

; <label>:11                                     ; preds = %5, %0
store i64 0, i64* %i, align 8
br label %12

```

The above code executes the `i = 0` statement of for loop.

```

; <label>:12                                     ; preds = %23, %11
%13 = load i64* %i, align 8
%14 = load i64* %numiter, align 8
%15 = icmp ult i64 344865, %14
br i1 %15, label %16, label %18

```

Label %12 executes the compare statement of for loop i.e. `i < MAX(MAGIC_NUMBER,numiter)`. %13 stores `i` from stack, %14 stores `numiter`, which is then compared with `MAGIC_NUMBER` which is 344865. It branches off depending on the compare result. `%15 = icmp ult i64 344865`, %14 is essentially executing the `MAX(MAGIC_NUMBER,numiter)` statement.

```

; <label>:16                                ; preds = %12
  %17 = load i64* %numiter, align 8
  br label %19

; <label>:18                                ; preds = %12
  br label %19

; <label>:19                                ; preds = %18, %16
  %20 = phi i64 [ %17, %16 ], [ 344865, %18 ]
  %21 = icmp ult i64 %13, %20
  br i1 %21, label %22, label %26

```

The above code is a continuation of the compare statement of for loop. Depending on the result of the `MAX()` function, %20 stores either `MAGIC_NUMBER` or `numiter`. `%21 = icmp ult i64 %13, %20` is the compare statement of for loop.

```

; <label>:22                                ; preds = %19
  br label %23

; <label>:23                                ; preds = %22
  %24 = load i64* %i, align 8
  %25 = add i64 %24, 1
  store i64 %25, i64* %i, align 8
  br label %12

```

Label %23 executes the update statement of for loop and then branches back to the beginning of the loop (label %12). `%25 = add i64 %24, 1` is the `i++` statement.

```

; <label>:26                                ; preds = %19
  ret i32 0

```

This is the final statement of the code which returns 0;

1.1.2 -O2

```

define i32 @emptyloop(i32 %argc, i8** nocapture readonly %argv) #0
{
    %1 = icmp sgt i32 %argc, 1
    br i1 %1, label %2, label %6

; <label>:2                                     ; preds = %0
    %3 = getelementptr inbounds i8** %argv, i64 1
    %4 = load i8** %3, align 8, !tbaa !1
    %5 = tail call i64 @strtol(i8* nocapture %4, i8** null, i32 10) #2
    br label %6

; <label>:6                                     ; preds = %2, %0
    ret i32 0

```

The above code shows the whole IR code using -O2 flag. We see that variables are not stored on stack but used directly as in the statement `%1 = icmp sgt i32 %argc, 1` where `%argc` is used directly. Also, the compiler has detected empty loop and omitted it from the IR. IR only consists of the if-statement. It's also using `strtol()` function instead of `atoi()` as mentioned in the source code. Using this, it doesn't require the sext instruction.

1.2 fib

1.2.1 -O0

```

define i32 @fib(i32 %n) #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    store i32 %n, i32* %2, align 4
    %3 = load i32* %2, align 4
    %4 = icmp slt i32 %3, 2
    br i1 %4, label %5, label %6

```

Initialising of variables on stack takes place. `%1` is used to store the return value. `%2` stores `n` on stack. `%4 = icmp slt i32 %3, 2` executes `n < 2` statement inside `if()`. It then branches depending on the result.

```

; <label>:5                                     ; preds = %0
    store i32 1, i32* %1
    br label %14

```

This stores 1 as the return value in case `n < 2`.

```

; <label>:6                                ; preds = %0
%7 = load i32* %2, align 4
%8 = sub nsw i32 %7, 1
%9 = call i32 @fib(i32 %8)
%10 = load i32* %2, align 4
%11 = sub nsw i32 %10, 2
%12 = call i32 @fib(i32 %11)
%13 = add nsw i32 %9, %12
store i32 %13, i32* %1
br label %14

```

This executes the recursive part of the program. %8 stores $n - 1$. %9 stores the result for $\text{fib}(n-1)$. Similarly, %12 stores $\text{fib}(n-2)$. The results are added in %13 and stored as the return value in %1 pointer.

```

; <label>:14                                ; preds = %6, %5
%15 = load i32* %1
ret i32 %15

```

This returns from the code using the return value stored at %1.

1.2.2 -O2

```

define i32 @fib(i32 %n) #0 {
  %1 = icmp slt i32 %n, 2
  br i1 %1, label %tailrecurse.\_crit\_edge, label
    %tailrecurse.preheader

```

No variable initialisation on stack. It creates new modules %tailrecurse, %tailrecurse._crit_edge etc. and branches to them directly based on the comparison result.

```

tailrecurse.preheader:                    ; preds = %0
  br label %tailrecurse

tailrecurse:                              ; preds =
  %tailrecurse, %tailrecurse.preheader
  %n.tr2 = phi i32 [ %4, %tailrecurse ], [ %n,
    %tailrecurse.preheader ]
  %accumulator.tr1 = phi i32 [ %5, %tailrecurse ], [ 1,
    %tailrecurse.preheader ]
  %2 = add nsw i32 %n.tr2, -1
  %3 = tail call i32 @fib(i32 %2)

```

```

%4 = add nsw i32 %n.tr2, -2
%5 = add nsw i32 %3, %accumulator.tr1
%6 = icmp slt i32 %4, 2
br i1 %6, label %tailrecurse._crit_edge.loopexit, label
    %tailrecurse

```

Here we can see the tail recursion optimization. The compiler uses tail call instead of call instruction for this. This avoids creating a new stack frame of the function (which happens in -O0) and uses a constant stack space instead. %accumulator.tr1 keeps track of the accumulated sum so far and %n.tr2 keeps track of the current value of n which keeps on updating in %4 = add nsw i32 %n.tr2, -2. For example, in case of factorial, the usual approach would be -

```

(fact 3)
(* 3 (fact 2))
(* 3 (* 2 (fact 1)))
(* 3 (* 2 (* 1 (fact 0))))
(* 3 (* 2 (* 1 1)))
(* 3 (* 2 1))
(* 3 2)
6

```

Whereas using tail-recursion, it would be -

```

(fact 3)
(fact-tail 3 1)
(fact-tail 2 3)
(fact-tail 1 6)
(fact-tail 0 6)
6

```

Here we keep track of the accumulated product so far (1, 3, 6, 6) and the current value of n (3, 2, 1, 0). phi instruction is used to keep track of the previous n.

```

tailrecurse._crit_edge.loopexit:                ; preds = %tailrecurse
    %.lcssa = phi i32 [ %5, %tailrecurse ]
    br label %tailrecurse._crit_edge

tailrecurse._crit_edge:                          ; preds =
    %tailrecurse._crit_edge.loopexit, %0
    %accumulator.tr.lcssa = phi i32 [ 1, %0 ], [ %.lcssa,
        %tailrecurse._crit_edge.loopexit ]

```

```
ret i32 %accumulator.tr.lcssa
```

loopexit is reached from the tailrecurse module and stores final result in `%.lcssa`. It then goes to `._crit_edge`. `._crit_edge` is reached either from loopexit or when the initial `n < 2` condition is failed. It gets the final result depending on the previous module using phi and finally returns it.

1.3 fibo_iter

1.3.1 -O0

```
define i64 @fibo_iter(i32 %n) #0 {  
  %1 = alloca i64, align 8  
  %2 = alloca i32, align 4  
  %fibo_cur = alloca i64, align 8  
  %fibo_prev = alloca i64, align 8  
  %i = alloca i32, align 4  
  %tmp = alloca i32, align 4  
  store i32 %n, i32* %2, align 4  
  %3 = load i32* %2, align 4  
  %4 = icmp ult i32 %3, 3  
  br i1 %4, label %5, label %6
```

Variables are allocated and stored on stack. `%1` stores the pointer to the result. `%2` stores pointer to `n`. `%4 = icmp ult i32 %3, 3` executes `n < 3`. The next `br` statement executes the if statement, branching depending on the result of compare (stored in `%4`).

```
; <label>:5                                ; preds = %0  
store i64 1, i64* %1  
br label %24
```

Stores 1 in the final return value (`%1` is the pointer to the final return value) in case `n < 3`.

```
; <label>:6                                ; preds = %0  
store i64 1, i64* %fibo_cur, align 8  
store i64 1, i64* %fibo_prev, align 8  
store i32 3, i32* %i, align 4  
br label %7
```

The above module is reached of `n <= 3`. It only initialises `fibo_cur`, `fibo_prev` and `i`.

```

; <label>:7                                ; preds = %19, %6
  %8 = load i32* %i, align 4
  %9 = load i32* %2, align 4
  %10 = icmp ule i32 %8, %9
  br i1 %10, label %11, label %22

```

%10 = icmp ule i32 %8, %9 statement executes the comparison statement of for loop (i.e. $i \leq n$). It then branches depending on the result.

```

; <label>:11                                ; preds = %7
  %12 = load i64* %fibonacci_cur, align 8
  %13 = trunc i64 %12 to i32
  store i32 %13, i32* %tmp, align 4
  %14 = load i64* %fibonacci_prev, align 8
  %15 = load i64* %fibonacci_cur, align 8
  %16 = add i64 %15, %14
  store i64 %16, i64* %fibonacci_cur, align 8
  %17 = load i32* %tmp, align 4
  %18 = zext i32 %17 to i64
  store i64 %18, i64* %fibonacci_prev, align 8
  br label %19

```

This is the main body of the loop. Note that `trunc` is used to typecast unsigned long to unsigned int and `zext`(zero extend) is used to typecast unsigned int to unsigned long. `store i32 %13, i32* %tmp, align 4` executes `tmp = fibonacci_cur`. %16 adds `fibonacci_cur` and `fibonacci_prev`. `store i64 %16, i64* %fibonacci_cur, align 8` stores the result of addition back in `fibonacci_cur`. The next 3 lines execute `fibonacci_prev = tmp`.

```

; <label>:19                                ; preds = %11
  %20 = load i32* %i, align 4
  %21 = add i32 %20, 1
  store i32 %21, i32* %i, align 4
  br label %7

```

This executes `i++` statement of for loop and then branches off to label 7 where the comparison statement of for loop takes place.

```

; <label>:22                                ; preds = %7
  %23 = load i64* %fibonacci_cur, align 8
  store i64 %23, i64* %1
  br label %24

```

```

; <label>:24                                     ; preds = %22, %5
  %25 = load i64* %1
  ret i64 %25

```

Label %22 stores the final result in the pointer pointing to return value (i.e. %1*). Label %24 returns the final value.

1.3.2 -O2

```

define i64 @fibonacci_iter(i32 %n) #0 {
  %1 = icmp ult i32 %n, 3
  br i1 %1, label %.loopexit, label %.lr.ph.preheader

```

This again creates new modules lr.ph, loopexit etc. for optimization. %1 = icmp ult i32 %n, 3 executes the if(n < 3) statement.

```

.lr.ph.preheader:                                ; preds = %0
  br label %.lr.ph

.lr.ph:                                           ; preds = %.lr.ph,
  %.lr.ph.preheader
  %i.03 = phi i32 [ %4, %.lr.ph ], [ 3, %.lr.ph.preheader ]
  %fibonacci_prev.02 = phi i64 [ %3, %.lr.ph ], [ 1, %.lr.ph.preheader ]
  %fibonacci_cur.01 = phi i64 [ %2, %.lr.ph ], [ 1, %.lr.ph.preheader ]
  %2 = add i64 %fibonacci_prev.02, %fibonacci_cur.01
  %3 = and i64 %fibonacci_cur.01, 4294967295
  %4 = add i32 %i.03, 1
  %5 = icmp ugt i32 %4, %n
  br i1 %5, label %.loopexit.loopexit, label %.lr.ph

```

The above code shows the main body of the loop. We see that the tmp variable has been omitted out. fibonacci_cur is being updated in the line %2 = add i64 %fibonacci_prev.02, %fibonacci_cur.01. To update fibonacci_prev, the compiler uses the number 4294967295 which is UINT_MAX. This is done because tmp and fibonacci_prev are of different types and hence integer overflow can occur. Adding UINT_MAX produces a result which consists of only those bits which tmp can store. The results are stored in temporary variables %2 and %3 and not updated directly. They are updated in the next iteration when these values are initialised in using phi instruction. Hence the value of fibonacci_cur in %3 = and i64 %fibonacci_cur.01, 4294967295 remains fibonacci_cur and is not fibonacci_cur + fibonacci_prev. If tmp were of type unsigned long, then we wouldn't have to add UINT_MAX and could've directly used fibonacci_cur.01 instead.

```

.loopexit.loopexit:                                ; preds = %.lr.ph
    %.lcssa = phi i64 [ %2, %.lr.ph ]
    br label %.loopexit

.loopexit:                                          ; preds =
    %.loopexit.loopexit, %0
    %.0 = phi i64 [ 1, %0 ], [ %.lcssa, %.loopexit.loopexit ]
    ret i64 %.0

```

This is the final part which returns the return value after exiting from the loop. `%.0 = phi i64 [1, %0], [%.lcssa, %.loopexit.loopexit]` chooses the return value based on whether we reached here after exiting the loop or directly after the initial if statement of `n < 3`.

1.4 gcd

1.4.1 -O0

```

define i32 @gcd1(i32 %a, i32 %b) #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    store i32 %a, i32* %2, align 4
    store i32 %b, i32* %3, align 4
    %4 = load i32* %3, align 4
    %5 = icmp ne i32 %4, 0
    br i1 %5, label %8, label %6

; <label>:6                                         ; preds = %0
    %7 = load i32* %2, align 4
    store i32 %7, i32* %1
    br label %14

; <label>:8                                         ; preds = %0
    %9 = load i32* %3, align 4
    %10 = load i32* %2, align 4
    %11 = load i32* %3, align 4
    %12 = srem i32 %10, %11
    %13 = call i32 @gcd1(i32 %9, i32 %12)
    store i32 %13, i32* %1
    br label %14

```

```

; <label>:14                                ; preds = %8, %6
  %15 = load i32* %1
  ret i32 %15

```

The above code defined the IR for gcd1(). Initially the variables are allocated on stack. %5 = icmp ne i32 %4, 0 executes !b statement. This and the next br statement implement the if statement. Label 6 is reached b is 0. It then stores a as the return value (in *%1). Otherwise it goes to label 8 where it recursively calls itself in the line %13 = call i32 @gcd1(i32 %9, i32 %12). Label 14 returns the final value.

```

define i32 @gcd2(i32 %a, i32 %b) #0 {
  %1 = alloca i32, align 4
  %2 = alloca i32, align 4
  store i32 %a, i32* %1, align 4
  store i32 %b, i32* %2, align 4
  br label %3

; <label>:3                                ; preds = %19, %0
  %4 = load i32* %1, align 4
  %5 = load i32* %2, align 4
  %6 = icmp ne i32 %4, %5
  br i1 %6, label %7, label %20

; <label>:7                                ; preds = %3
  %8 = load i32* %1, align 4
  %9 = load i32* %2, align 4
  %10 = icmp sgt i32 %8, %9
  br i1 %10, label %11, label %15

; <label>:11                               ; preds = %7
  %12 = load i32* %2, align 4
  %13 = load i32* %1, align 4
  %14 = sub nsw i32 %13, %12
  store i32 %14, i32* %1, align 4
  br label %19

; <label>:15                               ; preds = %7
  %16 = load i32* %1, align 4
  %17 = load i32* %2, align 4
  %18 = sub nsw i32 %17, %16
  store i32 %18, i32* %2, align 4
  br label %19

```

```

; <label>:19                                     ; preds = %15, %11
    br label %3

; <label>:20                                     ; preds = %3
    %21 = load i32* %1, align 4
    ret i32 %21

```

The above code shows the implementation for gcd2(). After storing a and b on stack, we branch to label 3 where we check whether a is equal to b or not in the line %6 = icmp ne i32 %4, %5. If they are equal, we branch off to label 20 which returns a. Otherwise we branch to label 7 where we then check if a < b or not in line %10 = icmp sgt i32 %8, %9. If a < b, we branch to label 11 where we execute a -= b. Similarly, we a >= b, we branch to label 15 where we execute b -= a. At the end of both label 11 and 15, we branch to label 19 and then to 3 which marks the beginning of the while loop.

```

define i32 @gcd3(i32 %a, i32 %b) #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %temp = alloca i32, align 4
    store i32 %a, i32* %1, align 4
    store i32 %b, i32* %2, align 4
    br label %3

; <label>:3                                     ; preds = %6, %0
    %4 = load i32* %2, align 4
    %5 = icmp ne i32 %4, 0
    br i1 %5, label %6, label %12

; <label>:6                                     ; preds = %3
    %7 = load i32* %2, align 4
    store i32 %7, i32* %temp, align 4
    %8 = load i32* %1, align 4
    %9 = load i32* %2, align 4
    %10 = srem i32 %8, %9
    store i32 %10, i32* %2, align 4
    %11 = load i32* %temp, align 4
    store i32 %11, i32* %1, align 4
    br label %3

; <label>:12                                     ; preds = %3

```

```

%13 = load i32* %1, align 4
ret i32 %13

```

The above code represents gcd3(). After initialising a, b, and temp on stack, we branch to label 3 where we check if b == 0 or not in line %5 = icmp ne i32 %4, 0. In case b == 0, we branch to label 12 where we return a. Else we branch to label 6 which is the main body of the function. %7 loads b and the next statement stores it in temp. Similarly, %8 and %9 stores a and b, and %10 stores a % b using %10 = srem i32 %8, %9. %11 stores temp and then stores in in a in the next statement. Finally we branch to label 3 which is the beginning of our while loop.

1.4.2 -O2

```

define i32 @gcd1(i32 %a, i32 %b) #0 {
    %1 = icmp eq i32 %b, 0
    br i1 %1, label %tailrecurse._crit_edge, label
        %tailrecurse.preheader

tailrecurse.preheader:                                ; preds = %0
    br label %tailrecurse

tailrecurse:                                           ; preds =
    %tailrecurse, %tailrecurse.preheader
    %b.tr2 = phi i32 [ %2, %tailrecurse ], [ %b,
        %tailrecurse.preheader ]
    %a.tr1 = phi i32 [ %b.tr2, %tailrecurse ], [ %a,
        %tailrecurse.preheader ]
    %2 = srem i32 %a.tr1, %b.tr2
    %3 = icmp eq i32 %2, 0
    br i1 %3, label %tailrecurse._crit_edge.loopexit, label
        %tailrecurse

tailrecurse._crit_edge.loopexit:                      ; preds = %tailrecurse
    %b.tr2.lcssa = phi i32 [ %b.tr2, %tailrecurse ]
    br label %tailrecurse._crit_edge

tailrecurse._crit_edge:                               ; preds =
    %tailrecurse._crit_edge.loopexit, %0
    %a.tr.lcssa = phi i32 [ %a, %0 ], [ %b.tr2.lcssa,
        %tailrecurse._crit_edge.loopexit ]
    ret i32 %a.tr.lcssa

```

The above code represents gcd1(). We check if b is zero or not in line %1 = icmp eq i32 %b, 0. If it's zero, we branch to tailrecurse._crit_edge and return a. Otherwise we branch to tailrecurse.preheader and then to tailrecurse. In tailrecurse, we get the values of a.tr1(first term of function) and b.tr2(second term of function) using phi instruction (depending on the previous module from which we entered it). %2 stores a % b which is then used as the second term in next iteration (i.e. b). If it's zero, (checked in %3 = icmp eq i32 %2, 0), then we exit the loop. Else we branch to tailrecurse again. After exiting the loop, we store the first term (i.e. a) as the return value in b.tr2.lcssa. The final return value is chosen depending on whether we reached there after recursing or directly after the first if loop (%a.tr.lcssa = phi i32 [%a, %0], [%b.tr2.lcssa, %tailrecurse._crit_edge.loopexit]).

```
define i32 @gcd2(i32 %a, i32 %b) #0 {
  %1 = icmp eq i32 %a, %b
  br i1 %1, label %.outer._crit_edge, label %.lr.ph.preheader

.lr.ph.preheader:                                     ; preds = %0
  br label %.lr.ph

.lr.ph:                                                ; preds = %.outer,
  %.lr.ph.preheader
  %.0.ph6 = phi i32 [ %.03.lcssa, %.outer ], [ %b,
  %.lr.ph.preheader ]
  %.01.ph5 = phi i32 [ %7, %.outer ], [ %a, %.lr.ph.preheader ]
  br label %4

; <label>:2                                           ; preds = %4
  %3 = icmp eq i32 %.01.ph5, %6
  br i1 %3, label %.outer._crit_edge.loopexit, label %4

; <label>:4                                           ; preds = %2, %.lr.ph
  %.03 = phi i32 [ %.0.ph6, %.lr.ph ], [ %6, %2 ]
  %5 = icmp slt i32 %.03, %.01.ph5
  %6 = sub nsw i32 %.03, %.01.ph5
  br i1 %5, label %.outer, label %2

.outer:                                              ; preds = %4
  %.03.lcssa = phi i32 [ %.03, %4 ]
  %7 = sub nsw i32 %.01.ph5, %.03.lcssa
  %8 = icmp eq i32 %7, %.03.lcssa
  br i1 %8, label %.outer._crit_edge.loopexit13, label %.lr.ph
```

```

.outer._crit_edge.loopexit:                                ; preds = %2
    %.01.ph5.lcssa = phi i32 [ %.01.ph5, %2 ]
    br label %.outer._crit_edge

.outer._crit_edge.loopexit13:                              ; preds = %.outer
    %.03.lcssa.lcssa = phi i32 [ %.03.lcssa, %.outer ]
    br label %.outer._crit_edge

.outer._crit_edge:                                         ; preds =
    %.outer._crit_edge.loopexit13, %.outer._crit_edge.loopexit, %0
    %.01.ph.lcssa = phi i32 [ %a, %0 ], [ %.01.ph5.lcssa,
        %.outer._crit_edge.loopexit ], [ %.03.lcssa.lcssa,
        %.outer._crit_edge.loopexit13 ]
    ret i32 %.01.ph.lcssa

```

The above code shows the IR for gcd2(). After comparing a and b in %1 = icmp eq i32 %a, %b, we branch to either outer._crit_edge (exit the loop and return a) or to lr.ph.header (enter the while loop). From lr.ph.header, we go to lr.ph which initialises 0.ph6 to b and 01.ph5 to a. Then we branch to label 4. Label 4 compares if a < b or not (%5 = icmp slt i32 %.03, %.01.ph5). If a < b, we branch to .outer where we store a - b in %7(%7 = sub nsw i32 %.01.ph5, %.03.lcssa). Since this will be our new a, we check here only if it's equal to b or not and exit the loop incase it is (%8 = icmp eq i32 %7, %.03.lcssa). If it's not, we branch to lr.ph again where the new a is initialised as %7 (which had a - b). Similarly, in case a >= b, we branch to label 2, where we check if a is equal to %6 (which stores b - a from .outer). If it is, we break from the loop and branch to outer._crit_edge. Else, we go to label 4 again where the new b is initialised as %6 (i.e. b - a). Since we exit the loop from many locations, we need to correctly initilise the return value. It'll either be the new value of a (i.e. a - b) stored in some temprary variable from previous module or it'll be a itself (in case the previous module had changed b i.e. executed b -= a). This is done using phi instruction. The final value is then returned.

1.5 loops

1.5.1 -O0

```

define zeroext i1 @is_sorted(i32* %a, i32 %n) #0 {
    %1 = alloca i1, align 1
    %2 = alloca i32*, align 8

```

```

%3 = alloca i32, align 4
%i = alloca i32, align 4
store i32* %a, i32** %2, align 8
store i32 %n, i32* %3, align 4
store i32 0, i32* %i, align 4
br label %4

; <label>:4                                ; preds = %24, %0
%5 = load i32* %i, align 4
%6 = load i32* %3, align 4
%7 = sub nsw i32 %6, 1
%8 = icmp slt i32 %5, %7
br i1 %8, label %9, label %27

; <label>:9                                ; preds = %4
%10 = load i32* %i, align 4
%11 = sext i32 %10 to i64
%12 = load i32** %2, align 8
%13 = getelementptr inbounds i32* %12, i64 %11
%14 = load i32* %13, align 4
%15 = load i32* %i, align 4
%16 = add nsw i32 %15, 1
%17 = sext i32 %16 to i64
%18 = load i32** %2, align 8
%19 = getelementptr inbounds i32* %18, i64 %17
%20 = load i32* %19, align 4
%21 = icmp sgt i32 %14, %20
br i1 %21, label %22, label %23

; <label>:22                                ; preds = %9
store i1 false, i1* %1
br label %28

; <label>:23                                ; preds = %9
br label %24

; <label>:24                                ; preds = %23
%25 = load i32* %i, align 4
%26 = add nsw i32 %25, 1
store i32 %26, i32* %i, align 4
br label %4

; <label>:27                                ; preds = %4

```



```

    store i1 true, i1* %1
    br label %28

; <label>:28                                ; preds = %27, %22
    %29 = load i1* %1
    ret i1 %29

```

The above code represents the `is_sorted()` function. Variables are initialized on stack. %1 stores the pointer to the return value, %2 stores pointer to `a[]`, %3 stores pointer to `n`. We then branch to label 4 after storing their values. Label 4 performs the comparison check of for loop. %7 stores `n - 1` and %5 stores `i`. They are compared in `%8 = icmp slt i32 %5, %7`. Depending on the result we branch to different labels. If `i ≤ n-1`, we branch to label 27 which stores `true` as the return value and then branches to label 28 which returns this value. Otherwise we branch to label 9 which is the inside of the loop. We load the stack values in temprary variables. %13 stores address of `a[i]`, %14 stores value of `a[i]`, %16 stores `i + 1`, %19 stores address of `a[i+1]`, %20 stores value of `a[i+1]`. We compare `a[i]` and `a[i+1]` using `%21 = icmp sgt i32 %14, %20`. If `a[i] < a[i + 1]`, we branch to label 22 which stores `false` as the return value and then branches to label 28 which returns this value. Otherwise we branch to label 23 and then to 24 where the update part of for loop takes place (i.e. `i++`). `%26 = add nsw i32 %25, 1` does `i++` and stores the result back in stack before branching to label 4 which is the beginning of the loop.

```

define void @add_arrays(i32* %a, i32* %b, i32* %c, i32 %n) #0 {
    %1 = alloca i32*, align 8
    %2 = alloca i32*, align 8
    %3 = alloca i32*, align 8
    %4 = alloca i32, align 4
    %i = alloca i32, align 4
    store i32* %a, i32** %1, align 8
    store i32* %b, i32** %2, align 8
    store i32* %c, i32** %3, align 8
    store i32 %n, i32* %4, align 4
    store i32 0, i32* %i, align 4
    br label %5

; <label>:5                                ; preds = %25, %0
    %6 = load i32* %i, align 4
    %7 = load i32* %4, align 4
    %8 = icmp slt i32 %6, %7

```

```

br i1 %8, label %9, label %28

; <label>:9                                ; preds = %5
%10 = load i32* %i, align 4
%11 = sext i32 %10 to i64
%12 = load i32** %1, align 8
%13 = getelementptr inbounds i32* %12, i64 %11
%14 = load i32* %13, align 4
%15 = load i32* %i, align 4
%16 = sext i32 %15 to i64
%17 = load i32** %2, align 8
%18 = getelementptr inbounds i32* %17, i64 %16
%19 = load i32* %18, align 4
%20 = add nsw i32 %14, %19
%21 = load i32* %i, align 4
%22 = sext i32 %21 to i64
%23 = load i32** %3, align 8
%24 = getelementptr inbounds i32* %23, i64 %22
store i32 %20, i32* %24, align 4
br label %25

; <label>:25                                ; preds = %9
%26 = load i32* %i, align 4
%27 = add nsw i32 %26, 1
store i32 %27, i32* %i, align 4
br label %5

; <label>:28                                ; preds = %5
ret void

```

The above represents the IR for `add_arrays` function. Again the variables are assigned space on stack and then initialised. %1, %2, %3 are pointers which stores address for arrays a, b, and c. %4 is a pointer to n. After initialising variables, we branch to label 5, where we compare i with n (%8 = `icmp slt i32 %6, %7`) i.e. the comparison statement of for loop. If $i \neq n$, we exit the loop by branching to label 28 which returns void. Otherwise we branch to label 9 which is the main body of the loop. Temporary variables are initialised again. %14 and %19 store the values of `a[i]` and `b[i]`. We store their sum in %20 (`%20 = add nsw i32 %14, %19`). %24 stores the address for `c[i]`. The sum is then stored at this location. (`store i32 %20, i32* %24, align 4`). We then branch to label 25 which updates the value of i i.e. `i++` (`%27 = add nsw i32 %26, 1`) and then stores it in stack (`store i32`

%27, i32* %i, align 4). We then branch to label 5 which is the beginning of the loop.

```

define i32 @sum(i8* %a, i32 %n) #0 {
    %1 = alloca i8*, align 8
    %2 = alloca i32, align 4
    %ret = alloca i32, align 4
    %i = alloca i32, align 4
    store i8* %a, i8** %1, align 8
    store i32 %n, i32* %2, align 4
    store i32 0, i32* %ret, align 4
    store i32 0, i32* %i, align 4
    br label %3

; <label>:3                                     ; preds = %16, %0
    %4 = load i32* %i, align 4
    %5 = load i32* %2, align 4
    %6 = icmp slt i32 %4, %5
    br i1 %6, label %7, label %19

; <label>:7                                     ; preds = %3
    %8 = load i32* %i, align 4
    %9 = sext i32 %8 to i64
    %10 = load i8** %1, align 8
    %11 = getelementptr inbounds i8* %10, i64 %9
    %12 = load i8* %11, align 1
    %13 = zext i8 %12 to i32
    %14 = load i32* %ret, align 4
    %15 = add nsw i32 %14, %13
    store i32 %15, i32* %ret, align 4
    br label %16

; <label>:16                                    ; preds = %7
    %17 = load i32* %i, align 4
    %18 = add nsw i32 %17, 1
    store i32 %18, i32* %i, align 4
    br label %3

; <label>:19                                    ; preds = %3
    %20 = load i32* %ret, align 4
    ret i32 %20

```

The above shows IR for sum() function. Variables are initialised on stack.

%1 is the pointer to array a, %2 is a pointer to n, %ret and %i are pointers to return value and i respectively. We branch to label 3 after initialization. In label 3, we perform the loop comparison operation i.e. `i < n` (`%6 = icmp slt i32 %4, %5`). If `i < n`, we exit the loop by branching to label 19 which returns the final sum. Otherwise we branch to label 7 which is the main body of the loop. We store addresses and values in temporary variables. %10 stores address of `a[i]`, %12 stores value of `a[i]`. `zext`(zero extends) is used to typecast. %14 stores the current value of `ret`. We add `a[i]` and `ret` using `%15 = add nsw i32 %14, %13` and store it back in stack (`store i32 %15, i32* %ret, align 4`). We then branch to label 16 where we update the value of `i`. %17 stores the current value of `i`. `%18 = add nsw i32 %17, 1` performs `i++` and `store i32 %18, i32* %i, align 4` stores it back in the stack. We branch back to label 3 which is the beginning of the loop.

```
define i32 @sumn(i32 %n) #0 {
    %1 = alloca i32, align 4
    %ret = alloca i32, align 4
    %i = alloca i32, align 4
    store i32 %n, i32* %1, align 4
    store i32 0, i32* %ret, align 4
    store i32 0, i32* %i, align 4
    br label %2

; <label>:2                                     ; preds = %0, %0
    %3 = load i32* %i, align 4
    %4 = load i32* %1, align 4
    %5 = icmp slt i32 %3, %4
    br i1 %5, label %6, label %13

; <label>:6                                     ; preds = %2
    %7 = load i32* %ret, align 4
    %8 = load i32* %i, align 4
    %9 = add nsw i32 %7, %8
    store i32 %9, i32* %ret, align 4
    br label %10

; <label>:10                                    ; preds = %6
    %11 = load i32* %i, align 4
    %12 = add nsw i32 %11, 1
    store i32 %12, i32* %i, align 4
    br label %2

; <label>:13                                    ; preds = %2
```

```
%14 = load i32* %ret, align 4
ret i32 %14
```

The above code shows the IR for `sumn()` function. The variables are initialised on stack. `%1` is pointer to `n`, `%ret` is pointer to `n`, `%i` is pointer to `i`. After initialising them on stack, we branch to label 2 where we check the comparison statement of for loop i.e. `if i < n` or `not(%5 = icmp slt i32 %3, %4)`. If `i <= n`, we branch to label 13 and return the return value. Otherwise we branch to label 6 which is the main body of the loop. `%7` stores the current value of `ret` and `%8` stores `i`. We add then (`%9 = add nsw i32 %7, %8`) and store it back in stack(`store i32 %9, i32* %ret, align 4`). We then branch to label 10 which updates the value of `i`(i.e. `i++`) before branching to label 2 again which is the beginning of the for loop.

1.5.2 -O2

```
define zeroext i1 @is_sorted(i32* nocapture readonly %a, i32 %n)
    #0 {
        %1 = add nsw i32 %n, -1
        %2 = sext i32 %1 to i64
        br label %3

; <label>:3                                     ; preds = %5, %0
        %indvars.iv = phi i64 [ %indvars.iv.next, %5 ], [ 0, %0 ]
        %4 = icmp slt i64 %indvars.iv, %2
        br i1 %4, label %5, label %11

; <label>:5                                     ; preds = %3
        %6 = getelementptr inbounds i32* %a, i64 %indvars.iv
        %7 = load i32* %6, align 4, !tbaa !1
        %indvars.iv.next = add nuw nsw i64 %indvars.iv, 1
        %8 = getelementptr inbounds i32* %a, i64 %indvars.iv.next
        %9 = load i32* %8, align 4, !tbaa !1
        %10 = icmp sgt i32 %7, %9
        br i1 %10, label %11, label %3

; <label>:11                                     ; preds = %5, %3
        %.0 = phi i1 [ false, %5 ], [ true, %3 ]
        ret i1 %.0
```

Above code shows IR for `is_sorted` function with `-O2` flag. `%1` stores `n - 1`. `%2` stores `n - 1` signed extended to unsigned long. We branch to label 3.

%indvars.iv stores the value of i , and %indvars.iv.next (in label 6) stores $i + 1$. %indvars.iv is initialised to 0 originally and then updated from label 6's value. We compare i with $n-1$ using `%4 = icmp slt i64 %indvars.iv, %2` and then branch accordingly. If $i \leq n-1$, we branch to label 5 which is the main body of the loop, else we branch to label 11 where we return the final value. In label 5, %6 stores address for $a[i]$, %7 stores value of $a[i]$, %indvars.iv.next stores $i + 1$, %8 stores address for $a[i + 1]$, %9 stores value of $a[i+1]$. We compare $a[i]$ and $a[i+1]$ using `%10 = icmp sgt i32 %7, %9`. If $a[i] \leq a[i + 1]$, we exit the loop by branching to label 11 and returning false. Else after the loop is done, we branch to label 11 and return true.

```
define void @add_arrays(i32* nocapture readonly %a, i32* nocapture
    readonly %b, i32* nocapture %c, i32 %n) #1 {
    %1 = icmp sgt i32 %n, 0
    br i1 %1, label %.lr.ph, label %._crit_edge

.lr.ph:                                     ; preds = %0
    %2 = add i32 %n, -1
    %3 = zext i32 %2 to i64
    %4 = add nuw nsw i64 %3, 1
    %end.idx = add nuw nsw i64 %3, 1
    %n.vec = and i64 %4, 8589934584
    %cmp.zero = icmp eq i64 %n.vec, 0
    %5 = add i32 %n, -1
    %6 = zext i32 %5 to i64
    %scevgep = getelementptr i32* %c, i64 %6
    br i1 %cmp.zero, label %middle.block, label %vector.memcheck

vector.memcheck:                           ; preds = %.lr.ph
    %scevgep5 = getelementptr i32* %a, i64 %6
    %bound0 = icmp uge i32* %scevgep5, %c
    %bound1 = icmp uge i32* %scevgep, %a
    %found.conflict = and i1 %bound0, %bound1
    %scevgep8 = getelementptr i32* %b, i64 %6
    %bound010 = icmp uge i32* %scevgep8, %c
    %bound111 = icmp uge i32* %scevgep, %b
    %found.conflict12 = and i1 %bound010, %bound111
    %conflict.rdx = or i1 %found.conflict, %found.conflict12
    br i1 %conflict.rdx, label %middle.block, label
        %vector.body.preheader

vector.body.preheader:                     ; preds =
    %vector.memcheck
```

```

br label %vector.body

vector.body:                                     ; preds =
    %vector.body, %vector.body.preheader
    %index = phi i64 [ %index.next, %vector.body ], [ 0,
        %vector.body.preheader ]
    %7 = getelementptr inbounds i32* %a, i64 %index
    %8 = bitcast i32* %7 to <4 x i32>*
    %wide.load = load <4 x i32>* %8, align 4, !tbaa !1
    %.sum23 = or i64 %index, 4
    %9 = getelementptr i32* %a, i64 %.sum23
    %10 = bitcast i32* %9 to <4 x i32>*
    %wide.load14 = load <4 x i32>* %10, align 4, !tbaa !1
    %11 = getelementptr inbounds i32* %b, i64 %index
    %12 = bitcast i32* %11 to <4 x i32>*
    %wide.load15 = load <4 x i32>* %12, align 4, !tbaa !1
    %.sum24 = or i64 %index, 4
    %13 = getelementptr i32* %b, i64 %.sum24
    %14 = bitcast i32* %13 to <4 x i32>*
    %wide.load16 = load <4 x i32>* %14, align 4, !tbaa !1
    %15 = add nsw <4 x i32> %wide.load15, %wide.load
    %16 = add nsw <4 x i32> %wide.load16, %wide.load14
    %17 = getelementptr inbounds i32* %c, i64 %index
    %18 = bitcast i32* %17 to <4 x i32>*
    store <4 x i32> %15, <4 x i32>* %18, align 4, !tbaa !1
    %.sum25 = or i64 %index, 4
    %19 = getelementptr i32* %c, i64 %.sum25
    %20 = bitcast i32* %19 to <4 x i32>*
    store <4 x i32> %16, <4 x i32>* %20, align 4, !tbaa !1
    %index.next = add i64 %index, 8
    %21 = icmp eq i64 %index.next, %n.vec
    br i1 %21, label %middle.block.loopexit, label %vector.body,
        !llvm.loop !5

middle.block.loopexit:                           ; preds = %vector.body
    br label %middle.block

middle.block:                                     ; preds =
    %middle.block.loopexit, %vector.memcheck, %.lr.ph
    %resume.val = phi i64 [ 0, %.lr.ph ], [ 0, %vector.memcheck ], [
        %n.vec, %middle.block.loopexit ]
    %cmp.n = icmp eq i64 %end.idx, %resume.val
    br i1 %cmp.n, label %._crit_edge, label %scalar.ph.preheader

```

```

scalar.ph.preheader:                                ; preds =
    %middle.block
    %22 = trunc i64 %resume.val to i32
    %23 = sub i32 %n, %22
    %24 = add i32 %n, -1
    %25 = sub i32 %24, %22
    %xtraiter = and i32 %23, 1
    %lcmp.mod = icmp ne i32 %xtraiter, 0
    br i1 %lcmp.mod, label %scalar.ph.prol, label
        %scalar.ph.preheader.split

scalar.ph.prol:                                      ; preds =
    %scalar.ph.preheader
    %26 = getelementptr inbounds i32* %a, i64 %resume.val
    %27 = load i32* %26, align 4, !tbaa !1
    %28 = getelementptr inbounds i32* %b, i64 %resume.val
    %29 = load i32* %28, align 4, !tbaa !1
    %30 = add nsw i32 %29, %27
    %31 = getelementptr inbounds i32* %c, i64 %resume.val
    store i32 %30, i32* %31, align 4, !tbaa !1
    %indvars.iv.next.prol = add nuw nsw i64 %resume.val, 1
    %lftr.wideiv.prol = trunc i64 %resume.val to i32
    %exitcond.prol = icmp eq i32 %lftr.wideiv.prol, %2
    br label %scalar.ph.preheader.split

scalar.ph.preheader.split:                          ; preds =
    %scalar.ph.prol, %scalar.ph.preheader
    %indvars.iv.unr = phi i64 [ %resume.val, %scalar.ph.preheader ],
        [ %indvars.iv.next.prol, %scalar.ph.prol ]
    %32 = icmp ult i32 %25, 1
    br i1 %32, label %._crit_edge.loopexit, label
        %scalar.ph.preheader.split.split

scalar.ph.preheader.split.split:                    ; preds =
    %scalar.ph.preheader.split
    br label %scalar.ph

scalar.ph:                                           ; preds = %scalar.ph,
    %scalar.ph.preheader.split.split
    %indvars.iv = phi i64 [ %indvars.iv.unr,
        %scalar.ph.preheader.split.split ], [ %indvars.iv.next.1,
        %scalar.ph ]

```



```

%33 = getelementptr inbounds i32* %a, i64 %indvars.iv
%34 = load i32* %33, align 4, !tbaa !1
%35 = getelementptr inbounds i32* %b, i64 %indvars.iv
%36 = load i32* %35, align 4, !tbaa !1
%37 = add nsw i32 %36, %34
%38 = getelementptr inbounds i32* %c, i64 %indvars.iv
store i32 %37, i32* %38, align 4, !tbaa !1
%indvars.iv.next = add nuw nsw i64 %indvars.iv, 1
%lftr.wideiv = trunc i64 %indvars.iv to i32
%39 = getelementptr inbounds i32* %a, i64 %indvars.iv.next
%40 = load i32* %39, align 4, !tbaa !1
%41 = getelementptr inbounds i32* %b, i64 %indvars.iv.next
%42 = load i32* %41, align 4, !tbaa !1
%43 = add nsw i32 %42, %40
%44 = getelementptr inbounds i32* %c, i64 %indvars.iv.next
store i32 %43, i32* %44, align 4, !tbaa !1
%indvars.iv.next.1 = add nuw nsw i64 %indvars.iv.next, 1
%lftr.wideiv.1 = trunc i64 %indvars.iv.next to i32
%exitcond.1 = icmp eq i32 %lftr.wideiv.1, %2
br i1 %exitcond.1, label %._crit_edge.loopexit.unr-lcssa, label
    %scalar.ph, !llvm.loop !8

._crit_edge.loopexit.unr-lcssa:                                ; preds = %scalar.ph
    br label %._crit_edge.loopexit

._crit_edge.loopexit:                                          ; preds =
    %._crit_edge.loopexit.unr-lcssa, %scalar.ph.preheader.split
    br label %._crit_edge

._crit_edge:                                                  ; preds =
    %._crit_edge.loopexit, %middle.block, %0
    ret void

```

The above code shows the `add_arrays` function. We initially check if `n < 0` or not. If not, we directly branch to `._crit_edge` and return void. Else we branch to label `._lr.ph`. Here, `%2` stores `n - 1`, and `%3` stores the same thing typecasted to unsigned long. `%4` and `%end.ind` stores `n`. `%n.vec` also stores `n` initially (8589934584 is the max possible value that can be stored in unsigned long). `%5` and `%6` are the same as `%2` and `%3` (duplicated possibly because we create separate variables for separate arrays). `%scevgep` stores the address for `c[i]`.

```

define i32 @sum(i8* nocapture readonly %a, i32 %n) #0 {

```

```

%1 = icmp sgt i32 %n, 0
br i1 %1, label %.lr.ph, label %._crit_edge

.lr.ph:                                     ; preds = %0
%2 = add i32 %n, -1
%3 = zext i32 %2 to i64
%4 = add nuw nsw i64 %3, 1
%end.idx = add nuw nsw i64 %3, 1
%n.vec = and i64 %4, 8589934584
%cmp.zero = icmp eq i64 %n.vec, 0
br i1 %cmp.zero, label %middle.block, label %vector.body.preheader

vector.body.preheader:                     ; preds = %.lr.ph
br label %vector.body

vector.body:                               ; preds =
    %vector.body, %vector.body.preheader
%index = phi i64 [ %index.next, %vector.body ], [ 0,
    %vector.body.preheader ]
%vec.phi = phi <4 x i32> [ %11, %vector.body ], [
    zeroinitializer, %vector.body.preheader ]
%vec.phi4 = phi <4 x i32> [ %12, %vector.body ], [
    zeroinitializer, %vector.body.preheader ]
%5 = getelementptr inbounds i8* %a, i64 %index
%6 = bitcast i8* %5 to <4 x i8>*
%wide.load = load <4 x i8>* %6, align 1, !tbaa !9
%.sum16 = or i64 %index, 4
%7 = getelementptr i8* %a, i64 %.sum16
%8 = bitcast i8* %7 to <4 x i8>*
%wide.load5 = load <4 x i8>* %8, align 1, !tbaa !9
%9 = zext <4 x i8> %wide.load to <4 x i32>
%10 = zext <4 x i8> %wide.load5 to <4 x i32>
%11 = add nsw <4 x i32> %9, %vec.phi
%12 = add nsw <4 x i32> %10, %vec.phi4
%index.next = add i64 %index, 8
%13 = icmp eq i64 %index.next, %n.vec
br i1 %13, label %middle.block.loopexit, label %vector.body,
    !llvm.loop !10

middle.block.loopexit:                     ; preds = %vector.body
%.lcssa18 = phi <4 x i32> [ %12, %vector.body ]
%.lcssa17 = phi <4 x i32> [ %11, %vector.body ]
br label %middle.block

```

```

middle.block:                                     ; preds =
    %middle.block.loopexit, %.lr.ph
    %resume.val = phi i64 [ 0, %.lr.ph ], [ %n.vec,
        %middle.block.loopexit ]
    %rdx.vec.exit.phi = phi <4 x i32> [ zeroinitializer, %.lr.ph ], [
        %.lcssa17, %middle.block.loopexit ]
    %rdx.vec.exit.phi12 = phi <4 x i32> [ zeroinitializer, %.lr.ph ],
        [ %.lcssa18, %middle.block.loopexit ]
    %bin.rdx = add <4 x i32> %rdx.vec.exit.phi12, %rdx.vec.exit.phi
    %rdx.shuf = shufflevector <4 x i32> %bin.rdx, <4 x i32> undef, <4
        x i32> <i32 2, i32 3, i32 undef, i32 undef>
    %bin.rdx13 = add <4 x i32> %bin.rdx, %rdx.shuf
    %rdx.shuf14 = shufflevector <4 x i32> %bin.rdx13, <4 x i32>
        undef, <4 x i32> <i32 1, i32 undef, i32 undef, i32 undef>
    %bin.rdx15 = add <4 x i32> %bin.rdx13, %rdx.shuf14
    %14 = extractelement <4 x i32> %bin.rdx15, i32 0
    %cmp.n = icmp eq i64 %end.idx, %resume.val
    br i1 %cmp.n, label %._crit_edge, label %scalar.ph.preheader

scalar.ph.preheader:                             ; preds =
    %middle.block
    %15 = trunc i64 %resume.val to i32
    %16 = sub i32 %n, %15
    %17 = add i32 %n, -1
    %18 = sub i32 %17, %15
    %xtraiter = and i32 %16, 3
    %lcmp.mod = icmp ne i32 %xtraiter, 0
    br i1 %lcmp.mod, label %scalar.ph.prol, label
        %scalar.ph.preheader.split

scalar.ph.prol:                                  ; preds =
    %scalar.ph.prol, %scalar.ph.preheader
    %indvars.iv.prol = phi i64 [ %indvars.iv.next.prol,
        %scalar.ph.prol ], [ %resume.val, %scalar.ph.preheader ]
    %ret.01.prol = phi i32 [ %22, %scalar.ph.prol ], [ %14,
        %scalar.ph.preheader ]
    %prol.iter = phi i32 [ %xtraiter, %scalar.ph.preheader ], [
        %prol.iter.sub, %scalar.ph.prol ]
    %19 = getelementptr inbounds i8* %a, i64 %indvars.iv.prol
    %20 = load i8* %19, align 1, !tbaa !9
    %21 = zext i8 %20 to i32
    %22 = add nsw i32 %21, %ret.01.prol

```

```

%indvars.iv.next.prol = add nuw nsw i64 %indvars.iv.prol, 1
%lftr.wideiv.prol = trunc i64 %indvars.iv.prol to i32
%exitcond.prol = icmp eq i32 %lftr.wideiv.prol, %2
%prol.iter.sub = sub i32 %prol.iter, 1
%prol.iter.cmp = icmp ne i32 %prol.iter.sub, 0
br i1 %prol.iter.cmp, label %scalar.ph.prol, label
    %scalar.ph.preheader.split, !llvm.loop !11

scalar.ph.preheader.split:                                ; preds =
    %scalar.ph.prol, %scalar.ph.preheader
    %.lcssa.unr = phi i32 [ 0, %scalar.ph.preheader ], [ %22,
        %scalar.ph.prol ]
    %indvars.iv.unr = phi i64 [ %resume.val, %scalar.ph.preheader ],
        [ %indvars.iv.next.prol, %scalar.ph.prol ]
    %ret.01.unr = phi i32 [ %14, %scalar.ph.preheader ], [ %22,
        %scalar.ph.prol ]
    %23 = icmp ult i32 %18, 3
    br i1 %23, label %._crit_edge.loopexit, label
        %scalar.ph.preheader.split.split

scalar.ph.preheader.split.split:                          ; preds =
    %scalar.ph.preheader.split
    br label %scalar.ph

scalar.ph:                                                ; preds = %scalar.ph,
    %scalar.ph.preheader.split.split
    %indvars.iv = phi i64 [ %indvars.iv.unr,
        %scalar.ph.preheader.split.split ], [ %indvars.iv.next.3,
        %scalar.ph ]
    %ret.01 = phi i32 [ %ret.01.unr, %scalar.ph.preheader.split.split
        ], [ %39, %scalar.ph ]
    %24 = getelementptr inbounds i8* %a, i64 %indvars.iv
    %25 = load i8* %24, align 1, !tbaa !9
    %26 = zext i8 %25 to i32
    %27 = add nsw i32 %26, %ret.01
    %indvars.iv.next = add nuw nsw i64 %indvars.iv, 1
    %lftr.wideiv = trunc i64 %indvars.iv to i32
    %28 = getelementptr inbounds i8* %a, i64 %indvars.iv.next
    %29 = load i8* %28, align 1, !tbaa !9
    %30 = zext i8 %29 to i32
    %31 = add nsw i32 %30, %27
    %indvars.iv.next.1 = add nuw nsw i64 %indvars.iv.next, 1
    %lftr.wideiv.1 = trunc i64 %indvars.iv.next to i32

```

```

%32 = getelementptr inbounds i8* %a, i64 %indvars.iv.next.1
%33 = load i8* %32, align 1, !tbaa !9
%34 = zext i8 %33 to i32
%35 = add nsw i32 %34, %31
%indvars.iv.next.2 = add nuw nsw i64 %indvars.iv.next.1, 1
%lftr.wideiv.2 = trunc i64 %indvars.iv.next.1 to i32
%36 = getelementptr inbounds i8* %a, i64 %indvars.iv.next.2
%37 = load i8* %36, align 1, !tbaa !9
%38 = zext i8 %37 to i32
%39 = add nsw i32 %38, %35
%indvars.iv.next.3 = add nuw nsw i64 %indvars.iv.next.2, 1
%lftr.wideiv.3 = trunc i64 %indvars.iv.next.2 to i32
%exitcond.3 = icmp eq i32 %lftr.wideiv.3, %2
br i1 %exitcond.3, label %._crit_edge.loopexit.unr-lcssa, label
    %scalar.ph, !llvm.loop !13

._crit_edge.loopexit.unr-lcssa:                                ; preds = %scalar.ph
    %lcssa.ph = phi i32 [ %39, %scalar.ph ]
    br label %._crit_edge.loopexit

._crit_edge.loopexit:                                         ; preds =
    %._crit_edge.loopexit.unr-lcssa, %scalar.ph.preheader.split
    %lcssa = phi i32 [ %lcssa.unr, %scalar.ph.preheader.split ], [
        %lcssa.ph, %._crit_edge.loopexit.unr-lcssa ]
    br label %._crit_edge

._crit_edge:                                                  ; preds =
    %._crit_edge.loopexit, %middle.block, %0
    %ret.0.lcssa = phi i32 [ 0, %0 ], [ %14, %middle.block ], [
        %lcssa, %._crit_edge.loopexit ]
    ret i32 %ret.0.lcssa

```

The above code shows the IR for `sum()` function.

```

define i32 @sumn(i32 %n) #2 {
    %1 = icmp sgt i32 %n, 0
    br i1 %1, label %lr.ph, label %11

lr.ph:                                                         ; preds = %0
    %2 = add i32 %n, -1
    %3 = zext i32 %2 to i33
    %4 = add i32 %n, -2
    %5 = zext i32 %4 to i33
    %6 = mul i33 %3, %5

```

```

%7 = lshr i33 %6, 1
%8 = trunc i33 %7 to i32
%9 = add i32 %8, %n
%10 = add i32 %9, -1
br label %11

; <label>:11                                ; preds = %.lr.ph, %0
%ret.0.lcssa = phi i32 [ %10, %.lr.ph ], [ 0, %0 ]
ret i32 %ret.0.lcssa

```

The above code shows IR for `sumn()` function. We initially check if `n` is zero or not. If yes, then we branch to label 11 and return 0. Else we branch to label `.lr.ph` which is the main body of the loop. Instead of iterating the loop, it calculates the return value using the formula for the sum of first `n` elements i.e.

$$\sum_{i=1}^n i = (n * (n + 1)) / 2$$

. `%3` stores `n - 1` (zero extended), and `%5` stores `n - 2` (zero extended). `%6` stored `(n - 1) * (n - 2)`. `%7 = lshr i33 %6, 1` shifts the product right by 1 bit essentially dividing it by 2. `%8` truncates it back to int from long. Till here, `%8` contains

$$((n - 1) * (n - 2)) / 2$$

which is the sum of first `n - 2` integers. But we want the sum of first `n - 1` integers. Hence we add `n - 1` to it using `%9 = add i32 %8, %n` and `%10 = add i32 %9, -1`. Finally we branch to label 11 and return the result.

1.6 print_arg

1.6.1 -O0

```

define i32 @print_arg(i32 %argc, i8** %argv) #0 {
  %1 = alloca i32, align 4
  %2 = alloca i32, align 4
  %3 = alloca i8**, align 8
  store i32 %argc, i32* %2, align 4
  store i8** %argv, i8*** %3, align 8
  %4 = load i32* %2, align 4
  %5 = icmp ne i32 %4, 2
  br i1 %5, label %6, label %7

; <label>:6                                ; preds = %0

```

```

store i32 -1, i32* %1
br label %12

; <label>:7                                     ; preds = %0
%8 = load i8*** %3, align 8
%9 = getelementptr inbounds i8** %8, i64 1
%10 = load i8** %9, align 8
%11 = call i32 @printf(i8** getelementptr inbounds ([3
    x i8]* @.str, i32 0, i32 0), i8* %10)
store i32 0, i32* %1
br label %12

; <label>:12                                     ; preds = %7, %6
%13 = load i32* %1
ret i32 %13

```

The above code shows the IR code for `print_arg` function. `%1` stores pointer to the return value. `%2` and `%3` stores pointers to `argc` and `argv[]`. We check `argc != 2` using `%5 = icmp ne i32 %4, 2`. If they're not equal, we branch to label 6 where the return value is initialised to -1. From there we branch to label 12 and return. Otherwise we branch to label 7 where the print statement takes place. We use `getelementptr` to find the address for `argv[1]` and load its value in `%10` (`%10 = load i8** %9, align 8`). We then call the print function using `%11 = call i32 @printf(i8** getelementptr inbounds ([3 x i8]* @.str, i32 0, i32 0), i8* %10)`. The return value is initialised to 1 before branching to label 12 where the return value is returned.

1.6.2 -O2

```

define i32 @print_arg(i32 %argc, i8** nocapture readonly %argv) #0
{
    %1 = icmp eq i32 %argc, 2
    br i1 %1, label %2, label %6

; <label>:2                                     ; preds = %0
%3 = getelementptr inbounds i8** %argv, i64 1
%4 = load i8** %3, align 8, !tbaa !1
%5 = tail call i32 @printf(i8** getelementptr inbounds
    ([3 x i8]* @.str, i64 0, i64 0), i8* %4) #2
br label %6

```

```

; <label>:6                                ; preds = %2, %0
%.0 = phi i32 [ 0, %2 ], [ -1, %0 ]
ret i32 %.0

```

The above code shows IR for `print_arg` using `-O2` optimization. We compare `argc` with 2 in `%1 = icmp eq i32 %argc, 2`. If they're equal, we branch to label 2. Otherwise we branch to label 6 from where we return. In label 2, `%3` stores the address for `argv[1]`, `%4` stores the value for `argv[1]`. Instead of calling `printf` directly, we use tail call `printf` for tail call optimization. Finally we branch to label 6. In label 6, the final return value is selected using `phi` instruction depending on the previous module from which we entered here (i.e. 0 if we entered from label 2 or -1 if we entered from label 0).

2 x86 Assembly

Similar to LLVM IR, it initialises stack space for variables and arguments. Tail call optimizations take place wherever possible and empty loops are discarded. For `sumn`, the sum is computed directly using closed formula. Return values may not be stored on stack and may be directly be stored in accumulator registers if possible.

2.1 emptyloop

2.1.1 -O0

```

emptyloop:                                # @emptyloop
    .cfi_startproc
# BB#0:
    pushl %ebp
.Ltmp0:
    .cfi_def_cfa_offset 8
.Ltmp1:
    .cfi_offset %ebp, -8
    movl %esp, %ebp
.Ltmp2:
    .cfi_def_cfa_register %ebp
    pushl %ebx
    pushl %esi
    subl $64, %esp
.Ltmp3:
    .cfi_offset %esi, -16

```



```

.Ltmp4:
    .cfi_offset %ebx, -12
    movl 12(%ebp), %eax
    movl 8(%ebp), %ecx
    movl %ecx, -12(%ebp)
    movl %eax, -16(%ebp)
    movl $0, -28(%ebp)
    movl $2147483646, -32(%ebp) # imm = 0x7FFFFFFE
    cmpl $2, -12(%ebp)
    jl .LBB0_2
# BB#1:
    movl -16(%ebp), %eax
    movl 4(%eax), %eax
    movl %esp, %ecx
    movl %eax, (%ecx)
    calll atoi
    movl %eax, %ecx
    sarl $31, %ecx
    movl %eax, -32(%ebp)
    movl %ecx, -28(%ebp)
.LBB0_2:
    movl $0, -20(%ebp)
    movl $0, -24(%ebp)
.LBB0_3:                                     # =>This Inner Loop Header:
    Depth=1
    movl -24(%ebp), %eax
    movl -20(%ebp), %ecx
    movl -32(%ebp), %edx
    movl -28(%ebp), %esi
    subl $344866, %edx                       # imm = 0x54322
    setb %bl
    movl %eax, -36(%ebp)                     # 4-byte Spill
    xorl %eax, %eax
    testl %esi, %esi
    movb %al, %bh
    movl %edx, -40(%ebp)                     # 4-byte Spill
    movl %ecx, -44(%ebp)                     # 4-byte Spill
    movb %bh, -45(%ebp)                     # 1-byte Spill
    movb %bl, -46(%ebp)                     # 1-byte Spill
    je .LBB0_11
# BB#10:                                     # in Loop: Header=BB0_3
    Depth=1
    movb -45(%ebp), %al                     # 1-byte Reload

```

```

    movb %al, -46(%ebp)      # 1-byte Spill
.LBB0_11:                    #   in Loop: Header=BB0_3
    Depth=1
    movb -46(%ebp), %al      # 1-byte Reload
    testb %al, %al
    jne .LBB0_5
    jmp .LBB0_4
.LBB0_4:                    #   in Loop: Header=BB0_3
    Depth=1
    movl -32(%ebp), %eax
    movl -28(%ebp), %ecx
    movl %eax, -52(%ebp)      # 4-byte Spill
    movl %ecx, -56(%ebp)      # 4-byte Spill
    jmp .LBB0_6
.LBB0_5:                    #   in Loop: Header=BB0_3
    Depth=1
    xorl %eax, %eax
    movl $344865, %ecx        # imm = 0x54321
    movl %ecx, -52(%ebp)      # 4-byte Spill
    movl %eax, -56(%ebp)      # 4-byte Spill
    jmp .LBB0_6
.LBB0_6:                    #   in Loop: Header=BB0_3
    Depth=1
    movl -56(%ebp), %eax      # 4-byte Reload
    movl -52(%ebp), %ecx      # 4-byte Reload
    movl -36(%ebp), %edx      # 4-byte Reload
    subl %ecx, %edx
    setae %bl
    movl -44(%ebp), %ecx      # 4-byte Reload
    subl %eax, %ecx
    setae %bh
    movl %ecx, -60(%ebp)      # 4-byte Spill
    movl %edx, -64(%ebp)      # 4-byte Spill
    movb %bh, -65(%ebp)       # 1-byte Spill
    movb %bl, -66(%ebp)       # 1-byte Spill
    je .LBB0_13
# BB#12:                    #   in Loop: Header=BB0_3
    Depth=1
    movb -65(%ebp), %al      # 1-byte Reload
    movb %al, -66(%ebp)       # 1-byte Spill
.LBB0_13:                    #   in Loop: Header=BB0_3
    Depth=1
    movb -66(%ebp), %al      # 1-byte Reload

```

```

    testb %al, %al
    jne .LBB0_9
    jmp .LBB0_7
.LBB0_7:                                #   in Loop: Header=BB0_3
    Depth=1
    jmp .LBB0_8
.LBB0_8:                                #   in Loop: Header=BB0_3
    Depth=1
    movl -24(%ebp), %eax
    movl -20(%ebp), %ecx
    addl $1, %eax
    adcl $0, %ecx
    movl %eax, -24(%ebp)
    movl %ecx, -20(%ebp)
    jmp .LBB0_3
.LBB0_9:
    xorl %eax, %eax
    addl $64, %esp
    popl %esi
    popl %ebx
    popl %ebp
    retl
.Ltmp5:
    .size emptyloop, .Ltmp5-emptyloop
    .cfi_endproc

```

2.1.2 -O2

```

emptyloop:                             # @emptyloop
    .cfi_startproc
# BB#0:
    subl $12, %esp
.Ltmp0:
    .cfi_def_cfa_offset 16
    cmpl $2, 16(%esp)
    jl .LBB0_2
# BB#1:
    movl 20(%esp), %eax
    movl 4(%eax), %eax
    movl %eax, (%esp)
    movl $10, 8(%esp)
    movl $0, 4(%esp)

```

```

    calll1 strtol
.LBB0_2:
    xorl %eax, %eax
    addl $12, %esp
    retl
.Ltmp1:
    .size emptyloop, .Ltmp1-emptyloop
    .cfi_endproc

```

The above code shows assembly program for emptyloop. It subtracts 12 from esp and sets the offset as 16. It compares argc with 2 using `cmpl $2, 16(%esp)`. In case `argc < 2`, it branches to .LBB0_2. There it initialises the return value to 0 using `xorl %eax, %eax` (eax stores the return value and xoring a register with itself makes it zero). It then restores the stack to its original state using `addl $12, %esp` before finally returning. In case `argc ≥ 2`, it gets the address of argc and stores in eax (`movl 20(%esp), %eax`). It stores `argc[1]` in eax using `movl 4(%eax), %eax`. Then it stores this value in the stack (`movl %eax, (%esp)`) as well as the other parameters that'll be required by the strtol function. It then calls the function (for implementing `atoi()`). (`calll1 strtol`). Finally it returns from the function. In other words, the loop portion has been omitted here also.

2.2 fib

2.2.1 -O0

```

fib:                                     # @fib
    .cfi_startproc
# BB#0:
    pushl %ebp
.Ltmp0:
    .cfi_def_cfa_offset 8
.Ltmp1:
    .cfi_offset %ebp, -8
    movl %esp, %ebp

```

It begins by pushing the current ebp on stack and defining offset. Then pushes the current esp in ebp to keep track the start of the current function stack.

```

.Ltmp2:
    .cfi_def_cfa_register %ebp

```

```

subl $24, %esp
movl 8(%ebp), %eax
movl %eax, -8(%ebp)
cmpl $2, -8(%ebp)
jge .LBB0_2

```

It compares n with 2 in line `cmpl $2, -8(%ebp)`. In case $n \geq 2$, it jumps to `.LBB0_2`.

```

# BB#1:
movl $1, -4(%ebp)
jmp .LBB0_3

```

If $n < 2$, it moves 1 to `ebp - 4` (i.e. stores 1 as the return value on stack), and jumps to `.LBB0_3`.

```

.LBB0_2:
movl -8(%ebp), %eax
subl $1, %eax
movl %eax, (%esp)
calll fib
movl -8(%ebp), %ecx
subl $2, %ecx
movl %ecx, (%esp)
movl %eax, -12(%ebp)      # 4-byte Spill
calll fib
movl -12(%ebp), %ecx      # 4-byte Reload
addl %eax, %ecx
movl %ecx, -4(%ebp)
.LBB0_3:
movl -4(%ebp), %eax
addl $24, %esp
popl %ebp
retl
.Ltmp3:
.size fib, .Ltmp3-fib
.cfi_endproc

```

In case $n \geq 2$, it stores n in `eax` (`movl -8(%ebp), %eax`) and subtracts 1 from it (`subl $1, %eax`). It moves it to stack and calls `fib` (which will now take the topmost value of stack i.e. $n-1$ as its argument). Similarly, it then stores $n-2$ in `ecx` and pushes it to stack. It then calls `fib` again (i.e. `fib(n-2)`). It adds the two results (`addl %eax, %ecx`) and moves it to `ebp - 4` (`movl %ecx,`

-4(%ebp)). This result is then stored in eax(which stores the return value) (movl -4(%ebp), %eax), stack is returned to it's original state and ebp is popped before returning.

The above code shows the assembly code for fib() function.

2.2.2 -O2

```

fib:                                     # @fib
    .cfi_startproc
# BB#0:
    pushl %edi
.Ltmp0:
    .cfi_def_cfa_offset 8
    pushl %esi
.Ltmp1:
    .cfi_def_cfa_offset 12
    pushl %eax
.Ltmp2:
    .cfi_def_cfa_offset 16
.Ltmp3:
    .cfi_offset %esi, -12
.Ltmp4:
    .cfi_offset %edi, -8
    movl 16(%esp), %edi
    movl $1, %esi
    cmpl $2, %edi
    jl .LBB0_2
    .align 16, 0x90
.LBB0_1:                                # %tailrecurse
                                        # =>This Inner Loop Header:
                                        Depth=1

    leal -1(%edi), %eax
    movl %eax, (%esp)
    calll fib
    addl $-2, %edi
    addl %eax, %esi
    cmpl $1, %edi
    jg .LBB0_1
.LBB0_2:                                # %tailrecurse._crit_edge
    movl %esi, %eax
    addl $4, %esp
    popl %esi
    popl %edi

```

```

    retl
.Ltmp5:
    .size fib, .Ltmp5-fib
    .cfi_endproc

```

Above code shows assembly code for `fib()` using `-O2` flag. After initialising stack and setting offsets, it moves `n` to `edi` (`movl 16(%esp), %edi`) and 1 to `edi` (which stores the return value in this case). It then compares `n` with 2 (`cmpl $2, %edi`). In case $n < 2$, it jumps to `.LBB0_2`. Else, it stores `n - 1` on stack and calls `fib`. Result is stored in `eax` (i.e. `eax` contains `fib(n-1)`). `n` becomes `n - 2` (`addl $-2, %edi`). Return value is updated (`addl %eax, %esi`). It compares `n - 2` with 1. In case $n - 2 > 1$ (or $n - 2 \geq 2$), it jumps back to the start of `.LBB0_1` to compute `fib(n-2)` (which it'll accumulate in `esi`). Finally, the return value is moved to `eax` in `.LBB0_2` (`movl %esi, %eax`). The stack is returned to it's former state before returning. In this case also we see tail recursion as it doesn't call `fib(n-2)` recursively, instead jumps back to an inner tail recursive loop.

2.3 fibo_iter

2.3.1 -O0

```

fibo_iter:                                # @fibo_iter
    .cfi_startproc
# BB#0:
    pushl %ebp
.Ltmp0:
    .cfi_def_cfa_offset 8
.Ltmp1:
    .cfi_offset %ebp, -8
    movl %esp, %ebp
.Ltmp2:
    .cfi_def_cfa_register %ebp
    pushl %esi
    subl $44, %esp
.Ltmp3:
    .cfi_offset %esi, -12
    movl 8(%ebp), %eax
    movl %eax, -20(%ebp)
    cmpl $3, -20(%ebp)
    jae .LBB0_2
# BB#1:

```

```

    movl $0, -12(%ebp)
    movl $1, -16(%ebp)
    jmp .LBB0_7
.LBB0_2:
    movl $0, -28(%ebp)
    movl $1, -32(%ebp)
    movl $0, -36(%ebp)
    movl $1, -40(%ebp)
    movl $3, -44(%ebp)
.LBB0_3:
    Depth=1
    movl -44(%ebp), %eax
    cmpl -20(%ebp), %eax
    ja .LBB0_6
# BB#4:
    Depth=1
    movl -32(%ebp), %eax
    movl %eax, -48(%ebp)
    movl -40(%ebp), %eax
    movl -36(%ebp), %ecx
    movl -32(%ebp), %edx
    movl -28(%ebp), %esi
    addl %eax, %edx
    adcl %ecx, %esi
    movl %edx, -32(%ebp)
    movl %esi, -28(%ebp)
    movl -48(%ebp), %eax
    movl %eax, -40(%ebp)
    movl $0, -36(%ebp)
# BB#5:
    Depth=1
    movl -44(%ebp), %eax
    addl $1, %eax
    movl %eax, -44(%ebp)
    jmp .LBB0_3
.LBB0_6:
    vmovsd -32(%ebp), %xmm0
    vmovsd %xmm0, -16(%ebp)
.LBB0_7:
    movl -16(%ebp), %eax
    movl -12(%ebp), %edx
    addl $44, %esp
    popl %esi

```

=>This Inner Loop Header:

in Loop: Header=BB0_3

in Loop: Header=BB0_3


```

    popl %ebp
    retl
.Ltmp4:
    .size fibo_iter, .Ltmp4-fibo_iter
    .cfi_endproc

```

The above code shows the assembly for `fibo_iter()` function.

2.3.2 -O2

```

fibo_iter:                                # @fibo_iter
    .cfi_startproc
# BB#0:
    pushl %ebp
.Ltmp0:
    .cfi_def_cfa_offset 8
    pushl %ebx
.Ltmp1:
    .cfi_def_cfa_offset 12
    pushl %edi
.Ltmp2:
    .cfi_def_cfa_offset 16
    pushl %esi
.Ltmp3:
    .cfi_def_cfa_offset 20
.Ltmp4:
    .cfi_offset %esi, -20
.Ltmp5:
    .cfi_offset %edi, -16
.Ltmp6:
    .cfi_offset %ebx, -12
.Ltmp7:
    .cfi_offset %ebp, -8
    movl 20(%esp), %ecx
    xorl %esi, %esi
    movl $1, %ebx
    cmpl $3, %ecx
    jae .LBB0_2
# BB#1:
    movl $1, %eax
    xorl %edx, %edx
    jmp .LBB0_4
.LBB0_2:

```

```

    movl $3, %edi
    movl $1, %ebp
    xorl %edx, %edx
    .align 16, 0x90
.LBB0_3:                                # %.lr.ph
                                        # =>This Inner Loop Header:
                                        Depth=1

    movl %ebx, %eax
    addl %ebp, %eax
    adcl %esi, %edx
    addl $1, %edi
    movl %ebp, %ebx
    movl %eax, %ebp
    cmpl %ecx, %edi
    jbe .LBB0_3
.LBB0_4:                                # %.loopexit
    popl %esi
    popl %edi
    popl %ebx
    popl %ebp
    retl
.Ltmp8:
    .size fibo_iter, .Ltmp8-fibo_iter
    .cfi_endproc

```

The above code shows the assembly for `fibo_iter()` using `-O2` flag.

2.4 print_arg

2.4.1 -O0

```

print_arg:                             # @print_arg
    .cfi_startproc
# BB#0:
    pushl %ebp
.Ltmp0:
    .cfi_def_cfa_offset 8
.Ltmp1:
    .cfi_offset %ebp, -8
    movl %esp, %ebp
.Ltmp2:
    .cfi_def_cfa_register %ebp
    subl $40, %esp

```

```

    movl 12(%ebp), %eax
    movl 8(%ebp), %ecx
    movl %ecx, -8(%ebp)
    movl %eax, -16(%ebp)
    cmpl $2, -8(%ebp)
    je .LBB0_2
# BB#1:
    movl $-1, -4(%ebp)
    jmp .LBB0_3
.LBB0_2:
    leal .L.str, %eax
    movl -16(%ebp), %ecx
    movl 4(%ecx), %ecx
    movl %eax, (%esp)
    movl %ecx, 4(%esp)
    calll printf
    movl $0, -4(%ebp)
    movl %eax, -20(%ebp)      # 4-byte Spill
.LBB0_3:
    movl -4(%ebp), %eax
    addl $40, %esp
    popl %ebp
    retl
.Ltmp3:
    .size print_arg, .Ltmp3-print_arg
    .cfi_endproc

```

Above code shows x86 assembly for `print_arg()` function. After moving the variables to stack and initialising it, it compares `n` with 2 (`cmpl $2, -8(%ebp)`). If equal, it jumps to `.LBB0_3` (where it restores the stack and sets the return value in `eax` before returning). Otherwise, it moves `argv[1]` to `ecx`, moves both of them to the top of stack using `movl %eax, (%esp)` and `movl %ecx, 4(%esp)` and calls the `printf` function. It changes the return value to 0 (`movl $0, -4(%ebp)`). Finally it returns as before (changing the stack to its original state etc.).

2.4.2 -O2

```

print_arg:                                # @print_arg
    .cfi_startproc
# BB#0:
    subl $12, %esp

```

```

.Ltmp0:
    .cfi_def_cfa_offset 16
    movl $-1, %eax
    cmpl $2, 16(%esp)
    jne .LBB0_2
# BB#1:
    movl 20(%esp), %eax
    movl 4(%eax), %eax
    movl %eax, 4(%esp)
    movl $.L.str, (%esp)
    calll printf
    xorl %eax, %eax
.LBB0_2:
    addl $12, %esp
    retl
.Ltmp1:
    .size print_arg, .Ltmp1-print_arg
    .cfi_endproc

```

Above code shows x86 assembly for `print_arg()` function using `-O2` flag. After initialising stack and setting offsets, it stores `-1` as the return value in `eax` (`movl $-1, %eax`). It compares `2` to `argc` in line `cmpl $2, 16(%esp)`. In case `2 != argc`, it jumps to `.LBB0_2` (where it simply restores the stack and returns). If `argc == 2`, it stores value of `argv[0]` in `eax` (`movl 20(%esp), %eax`), then stores value of `argv[1]` in `eax` (`movl 4(%eax), %eax`). It moves this value to the top of the stack. Then moves `.L.str` to the top of the stack (`.L.str` contains `%s`). It then calls the `printf` function. Finally it makes the return value as `0` (`xorl %eax, %eax`) and then restores the stack to its original state before returning.

3 Performance comparison across optimization levels

The following graphs show comparison between different optimization levels in llvm bitcode -

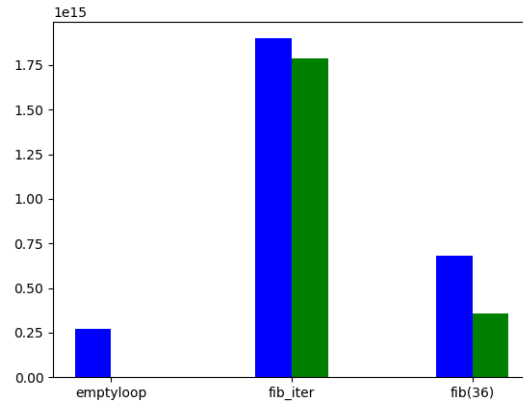


Figure 1: Comparison between different llvm optimizations

We see that -O3 performs slightly better than it's -O2 counterpart.

The following graph show comparison between different optimization levels in x86 assembly code -




Figure 2: Comparison between different x86 assembly optimizations

We see that there's not uch difference between -O2 and -O3 levels with -O2 sometimes performing better than -O3 (e.g. in fib_iter). However, there's a huge difference between -O0 and -O2/-O3.